


## ▼ Importing libraries

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
```

## ▼ Reading data

```
dataset = pd.read_csv('Powerplant_Data.csv')
# print(dataset.isnull().values.any()) - gave false, so no missing values
print(dataset.head())
x = dataset.iloc[:, :-1].values
y = dataset.iloc[:, -1].values
```



	AT	V	AP	RH	PE
0	14.96	41.76	1024.07	73.17	463.26
1	25.18	62.96	1020.04	59.08	444.37
2	5.11	39.40	1012.16	92.14	488.56
3	20.86	57.32	1010.24	76.64	446.48
4	10.82	37.50	1009.23	96.62	473.90

- Ambient Pressure (AP) in the range 992.89-1033.30 milibar,
- Relative Humidity (RH) in the range 25.56% to 100.16%
- Exhaust Vacuum (V) in teh range 25.36-81.56 cm Hg
- Net hourly electrical energy output (EP)

we are trying to predict EP based on the other features there is no categorical data so we don't need to do any encoding

## ▼ Splitting the data set for training and testing

```
from sklearn.model_selection import train_test_split
x_train, x_test, y_train, y_test = train_test_split(x, y, test_size=0.2, random_state=0)
```

## ▼ Multiple Regression

```
from sklearn.linear_model import LinearRegression
lin_reg = LinearRegression()
lin_reg.fit(x_train, y_train)

LinearRegression(copy_X=True, fit_intercept=True, n_jobs=None,
                 normalize=False)
```

## ▼ Polynomial Regression

```
from sklearn.preprocessing import PolynomialFeatures
poly_reg = PolynomialFeatures(degree = 4)
x_poly_train = poly_reg.fit_transform(x_train)
lin_reg_2 = LinearRegression()
lin_reg_2.fit(x_poly_train, y_train)

LinearRegression(copy_X=True, fit_intercept=True, n_jobs=None,
                 normalize=False)
```

## ▼ SVR

```
from sklearn.svm import SVR
# have to standardize x_train and y_train
from sklearn.preprocessing import StandardScaler
sc_x = StandardScaler()
sc_y = StandardScaler()
x_train_std = sc_x.fit_transform(x_train)
# y has to be a matrix
y_train_std = y_train.reshape(len(y_train), 1)
y_train_std = sc_y.fit_transform(y_train_std)

svr = SVR(kernel='rbf')
svr.fit(x_train_std, y_train_std)

C:\ProgramData\Anaconda3\lib\site-packages\sklearn\utils\validation.py:761: DataConversionWarning:
  y = column_or_1d(y, warn=True)
SVR(C=1.0, cache_size=200, coef0=0.0, degree=3, epsilon=0.1,
    gamma='auto_deprecated', kernel='rbf', max_iter=-1, shrinking=True,
    tol=0.001, verbose=False)
```

## ▼ Decision Tree

```
from sklearn.tree import DecisionTreeRegressor
tree_reg = DecisionTreeRegressor(random state=0)
```

```

tree_reg.fit(x_train, y_train)

DecisionTreeRegressor(criterion='mse', max_depth=None, max_features=None,
                      max_leaf_nodes=None, min_impurity_decrease=0.0,
                      min_impurity_split=None, min_samples_leaf=1,
                      min_samples_split=2, min_weight_fraction_leaf=0.0,
                      presort=False, random_state=0, splitter='best')

```

## ▼ Random Forest

```

from sklearn.ensemble import RandomForestRegressor
forest_reg = RandomForestRegressor(n_estimators=10, random_state=0)
forest_reg.fit(x_train, y_train)

RandomForestRegressor(bootstrap=True, criterion='mse', max_depth=None,
                      max_features='auto', max_leaf_nodes=None,
                      min_impurity_decrease=0.0, min_impurity_split=None,
                      min_samples_leaf=1, min_samples_split=2,
                      min_weight_fraction_leaf=0.0, n_estimators=10, n_jobs=None,
                      oob_score=False, random_state=0, verbose=0, warm_start=False)

```

## ▼ Model evaluation

```

from sklearn import metrics
# multiple linear regression
lin_reg_score = metrics.r2_score(y_test, lin_reg.predict(x_test))
print("Multiple Linear regression score: ", lin_reg_score)

# random forest regression
forest_reg_score = metrics.r2_score(y_test, forest_reg.predict(x_test))
# from 10 to 500, it didn't improve much
#TODO: how do you set the number of instances
print("Random Forest regression score: ", forest_reg_score)

#decision regression
tree_reg_score = metrics.r2_score(y_test, tree_reg.predict(x_test))
print("Decision Tree regression score: ", tree_reg_score)

# random forest regression
x_poly_test = poly_reg.fit_transform(x_test)
poly_reg_score = metrics.r2_score(y_test, lin_reg_2.predict(x_poly_test))
print("Polynomial regression score: ", poly_reg_score)

# svr
y_pred_svr = sc_y.inverse_transform(svr.predict(sc_x.transform(x_test)))
svr_score = metrics.r2_score(y_test, y_pred_svr)
print("SVR score: ", svr_score)

```

```
print( svr_score: , svr_score)
```

```
Multiple Linear regression score: 0.9325315554761303  
Random Forest regression score: 0.9615980699813017  
Decision Tree regression score: 0.9226091050550043  
Polynomial regression score: 0.945819341122773  
SVR score: 0.9480784049986264
```